

Threading, Parallelism and Networking

Slides inspired from [here](#) and the awesome book by Remzi

Plan

- Why threads
- Basics of Pthreads API
- Introducing Locks, Conditional variables and Semaphores
- A common deadlock scenario
- Basics of networking and SOCKET API
- If time permits: discussing a popular synchronization problem

Process vs. threads

- Parent P forks a child C
 - P and C do not share any memory
 - Need complicated IPC mechanisms to communicate
 - Extra copies of code, data in memory
- Parent P executes two threads T1 and T2
 - T1 and T2 share parts of the address space
 - Global variables can be used for communication
 - Smaller memory footprint
- Threads are like separate processes, except they share the same address space




Figure 26.1: Single-Threaded And Multi-Threaded Address Spaces

Why threads?

- Parallelism: a single process can effectively utilize multiple CPU cores
 - Understand the difference between concurrency and parallelism
 - **Concurrency**: running multiple threads/processes at the same time, **even on single CPU core**, by interleaving their executions
 - **Parallelism**: running multiple threads/processes in parallel over **different CPU cores**
- Even if no parallelism, concurrency of threads ensures effective use of CPU when one of the threads blocks (e.g., for I/O)

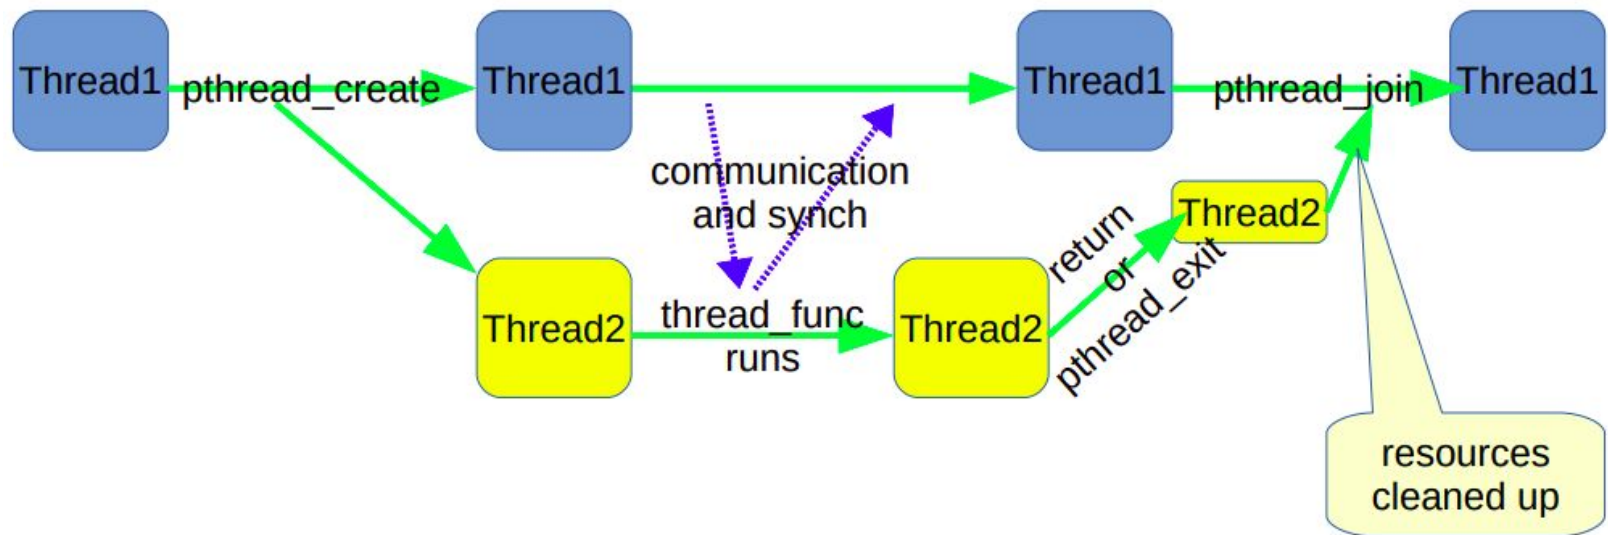
Creating a Thread

- `int pthread_create(pthread_t *tID, const pthread_attr_t *attr, void *(*thread_func)(void*), void *arg);`
- Create a new thread that will execute `thread_func`, passing it `arg`.
 - To pass the function multiple arguments, put them in a struct and pass a pointer to the struct
 - `*attr` - set thread attributes, `NULL` for default settings
- returns 0 if successful, an error number otherwise
 - updates `tID` with the thread ID of the new thread if successful

Creating threads using pthreads API

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4
5  void *mythread(void *arg) {
6      printf("%s\n", (char *) arg);
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     pthread_t p1, p2;
13     int rc;
14     printf("main: begin\n");
15     rc = pthread_create(&p1, NULL, mythread, "A"); assert(rc == 0);
16     rc = pthread_create(&p2, NULL, mythread, "B"); assert(rc == 0);
17     // join waits for the threads to finish
18     rc = pthread_join(p1, NULL); assert(rc == 0);
19     rc = pthread_join(p2, NULL); assert(rc == 0);
20     printf("main: end\n");
21     return 0;
22 }
```

Figure 26.2: Simple Thread Creation Code (t0.c)



Threads incrementing a counter without mutual exclusion

```
13
14 void *inc_counter(void *inp)
15 {
16     int thread_idx = ((struct thread_details *)inp)->idx;
17
18     // pthread_mutex_lock(&cnt_lock);
19     counter++;
20     // pthread_mutex_unlock(&cnt_lock);
21     return NULL;
22 }
23
24 int main()
25 {
26
27     const int TIMES = 1000;
28     pthread_t thread_ids_arr[TIMES];
29
30     pthread_mutex_init(&cnt_lock, NULL);
31
32     for (int i = 0; i < TIMES; i++)
33     {
34         pthread_t curr_tid;
35         td *thread_input = (td *) (malloc(sizeof(td)));
36         thread_input->idx = i;
37         pthread_create(&curr_tid, NULL, inc_counter, (void *) (thread_input));
38         thread_ids_arr[i] = curr_tid;
39     }
40     for (int i = 0; i < TIMES; i++)
41     {
42         pthread_join(thread_ids_arr[i], NULL);
43     }
44
45     pthread_mutex_destroy(&cnt_lock);
46     printf("Value of counter is: %d\n", counter);
47     return 0;
48 }
49
```


Why is this happening ?

OS	Thread 1	Thread 2	(after instruction)		
			PC	eax	counter
	<i>before critical section</i>		100	0	50
	mov 8049a1c,%eax		105	50	50
	add \$0x1,%eax		108	51	50
interrupt					
	<i>save T1</i>				
	<i>restore T2</i>		100	0	50
		mov 8049a1c,%eax	105	50	50
		add \$0x1,%eax	108	51	50
		mov %eax,8049a1c	113	51	51
interrupt					
	<i>save T2</i>				
	<i>restore T1</i>		108	51	51
	mov %eax,8049a1c		113	51	51

Figure 26.7: The Problem: Up Close and Personal

Locks: Basic idea

- Consider update of shared variable
counter = counter + 1
- We can use a special lock variable to protect it

```
1  lock_t mutex; // some globally-allocated lock 'mutex'  
2  ...  
3  lock(&mutex);  
4  balance = balance + 1;  
5  unlock(&mutex);
```

- All threads accessing a critical section share a lock
- One thread succeeds in locking – owner of lock
- Other threads that try to lock cannot proceed further until lock is released by the owner
- Pthreads library in Linux provides such locks

Syntax for locks

- **Declaration**

- `pthread_mutex_t lock;`

- **Initialization**

- `pthread_mutex_init(&lock, NULL)`

- **Locking the mutex**

- `pthread_mutex_lock(&lock)`

- **Unlocking the mutex**

- `pthread_mutex_unlock(&lock);`

- **Destroying the mutex**

- `pthread_mutex_destroy(&lock);`

Another type of synchronization

- Locks allow one type of synchronization between threads - mutual exclusion
- Another common requirement in multi-threaded applications - waiting and signalling
 - Example: Thread T1 wants to continue only after Thread T2 has finished some task
- Can accomplish this task by busy waiting on some variable - Inefficient
- Need a new synchronisation primitive - Condition variables

Condition Variables

- A condition variable (CV) is a queue that a thread can put itself into when waiting on some condition
- Another thread that makes the condition true can signal the CV to wake up the waiting thread
- Pthreads provides CV for user programs
 - OS has a similar functionality of wait/signal for kernel threads
- Signal wakes up one thread, signal broadcast wakes up all waiting threads

Parent waiting for Child - Condition Variables

```
1  int done  = 0;
2  pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3  pthread_cond_t c  = PTHREAD_COND_INITIALIZER;
4
5  void thr_exit() {
6      Pthread_mutex_lock(&m);
7      done = 1;
8      Pthread_cond_signal(&c);
9      Pthread_mutex_unlock(&m);
10 }
11
12 void *child(void *arg) {
13     printf("child\n");
14     thr_exit();
15     return NULL;
16 }
17
18 void thr_join() {
19     Pthread_mutex_lock(&m);
20     while (done == 0)
21         Pthread_cond_wait(&c, &m);
22     Pthread_mutex_unlock(&m);
23 }
24
25 int main(int argc, char *argv[]) {
26     printf("parent: begin\n");
27     pthread_t p;
28     Pthread_create(&p, NULL, child, NULL);
29     thr_join();
30     printf("parent: end\n");
31     return 0;
32 }
```

Why use lock when calling wait?


What if no lock is held when calling wait/signal?

- Race condition : missed wakeup
 - Parent checks done to be 0, decides to sleep, interrupted
 - Child runs, sets done to 1, signals, but no one sleeping yet
 - Parent now resumes and goes to sleep forever
- Lock must be held while calling wait and signal with CV
- The wait function releases the lock before putting thread to sleep, so lock is available for signaling thread

What is a semaphore?

- Synchronization primitive like conditional variables
- Semaphore is a variable with an underlying counter
- Two functions on a semaphore variable
 - Up/post increments the counter
 - Down/wait decrements the counter and blocks the calling thread if the resulting value is negative
- A semaphore with init value 1 acts as a simple lock
- **What does the current value of a semaphore mean?**
 - Positive : Number of threads that can decrement without blocking
 - Negative : Number of threads that have been blocked and are waiting
 - Zero : No threads are waiting, but if a thread tries to decrement it will block

Producer-Consumer Problem



- Producer waits for empty slot, inserts item in buffer, and notifies consumer
- Consumer waits for item, removes it from buffer, and notifies producer

Solution using Semaphores

- Need two semaphores for signaling
 - One to track empty slots, and make producer wait if no more empty slots
 - One to track filled slots, and make consumer wait if no more filled slots
- One semaphore to act as mutex for buffer

Incorrect, why?

Initial Values

- empty - Capacity of buffer
- Full - 0

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex);           // Line P0 (NEW LINE)
        sem_wait(&empty);           // Line P1
        put(i);                     // Line P2
        sem_post(&full);             // Line P3
        sem_post(&mutex);           // Line P4 (NEW LINE)
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex);           // Line C0 (NEW LINE)
        sem_wait(&full);            // Line C1
        int tmp = get();            // Line C2
        sem_post(&empty);           // Line C3
        sem_post(&mutex);           // Line C4 (NEW LINE)
        printf("%d\n", tmp);
    }
}
```

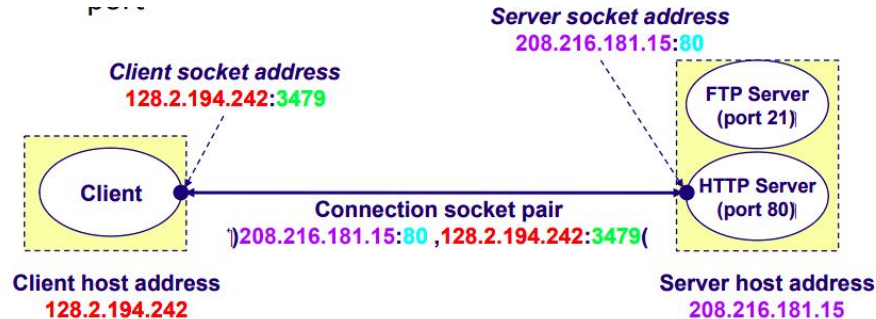
Correct Solution

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&empty);           // Line P1
        sem_wait(&mutex);           // Line P1.5 (MUTEX HERE)
        put(i);                     // Line P2
        sem_post(&mutex);           // Line P2.5 (AND HERE)
        sem_post(&full);            // Line P3
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&full);            // Line C1
        sem_wait(&mutex);           // Line C1.5 (MUTEX HERE)
        int tmp = get();            // Line C2
        sem_post(&mutex);           // Line C2.5 (AND HERE)
        sem_post(&empty);           // Line C3
        printf("%d\n", tmp);
    }
}
```

Sockets

- Socket API lets two processes in different machines to communicate with each other over the TCP/IP network stack
- Think of socket as the **door (logical pipe)** a process/thread needs to use if it needs to communicate with other processes
- **TCP sockets:** reliable delivery, congestion control
- Network socket identified by a port number on a machine – Socket bound to IP address of a network interface and a port number
- TCP socket communication: a “server” listens on a well-known port number, a “client” connects to the server, both exchange messages



Simple case when there is just one thread on the server side

TCP client

- socket: create new socket "C1" (returns socket file descriptor)
- connect: connects to a server socket using the server's IP address and port number (initiates TCP 3 way handshake to server). This request is intercepted by server socket 'A' and if connection is successful, socket C1 can now be used to communicate with the server
- read: read data from a connected socket "C1" [receives what server writes on its own socket "C2"]
- write: write data from a memory buffer into socket "C1" (which server receives on "C2")

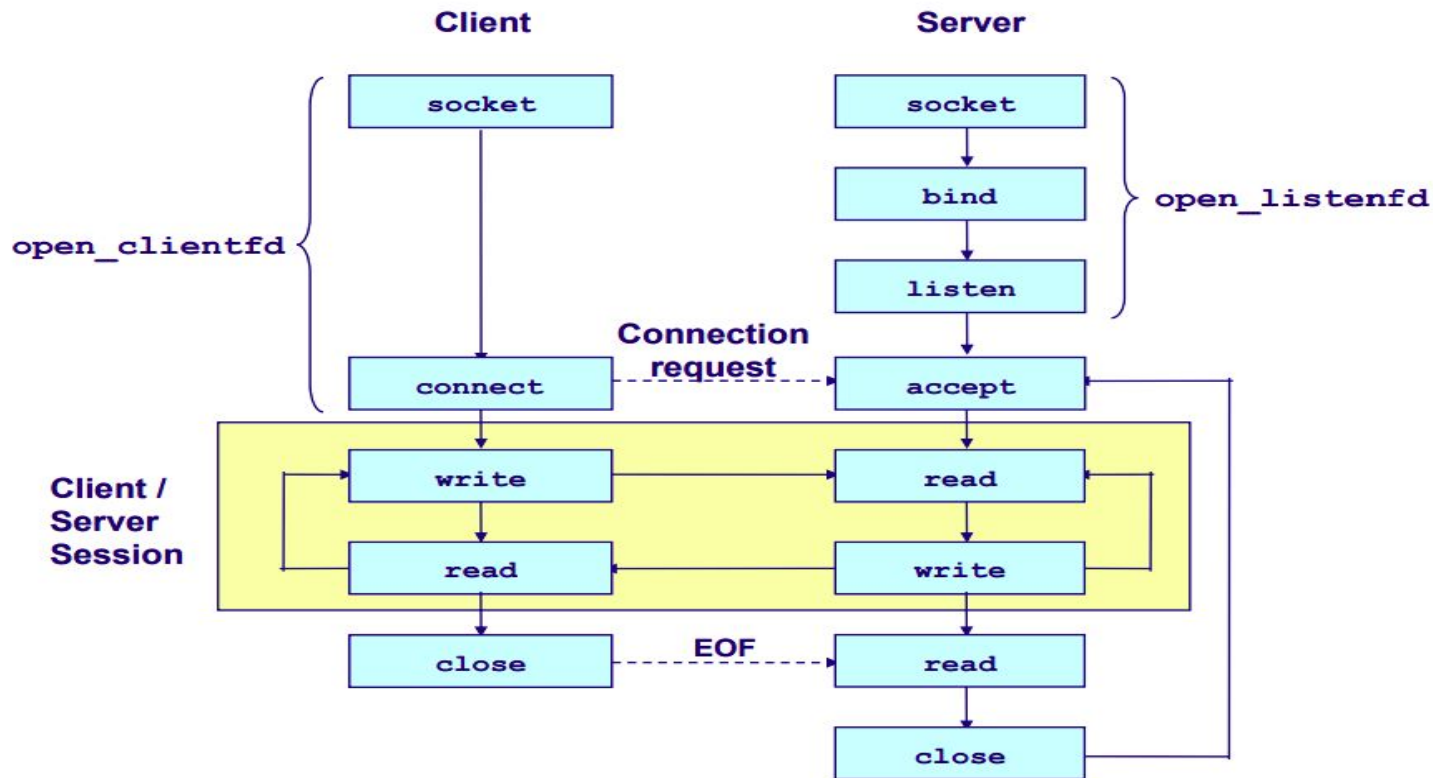
TCP server

In the simple case, the server has 2 sockets:

1. Listening socket (A) where server listens for connection requests
 2. Communication socket (denoted by C2 below) which is used to send and receive messages
- socket: create new socket "A" (returns socket file descriptor)
 - bind: bind server socket "A" to well-known port number and IP address
 - listen: start listening for new connections on socket "A"
 - accept: accept a new connection on server socket "A" (returns a new socket "C2" to talk to this specific client)
 - read: read data from socket "C2" into memory buffer [receives what client writes on its own socket "C1"]
 - write: write data from memory buffer into socket "C" connected to client (which then client receives on "C1")

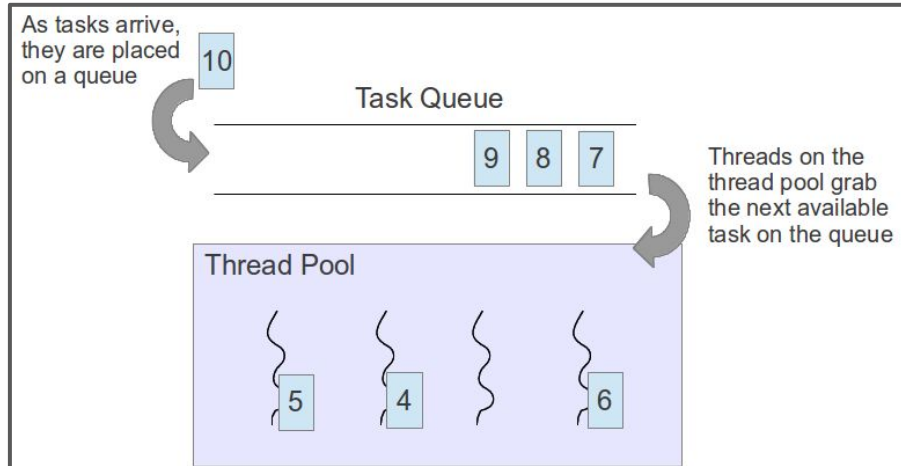

A TCP/UDP connection is identified by a tuple of five values:

{<protocol>, <src addr>, <src port>, <dest addr>, <dest port>}



Few components of the Socket API

- `int s = socket(domain, type, protocol) ;` - Create a socket
 - domain: Communication domain, typically used `AF_INET` (IPv4 protocol)
 - type: Type of socket - `SOCK_STREAM` (TCP) or `SOCK_DGRAM` (UDP)
 - protocol : Specifies protocols - usually set to 0 [Explore !]
- `int success_status = bind(sockid, &addrport, size);` - Reserves a port for the socket
 - sockid: socket identifier
 - addrport: struct `sockaddr_in` - the (IP) address and port of the machine (address usually set to `INADDR_ANY` chooses a local address)
 - size : Size of the `sockaddr` structure
- About struct `socladdr_in`:
 - **sin_family** : Set this to `AF_INET` ; (used to designate the type of addresses that your socket can communicate with)
 - **sin_port** : The network byte-ordered bit port number
 - **sin_addr** : Source address, `INADDR_ANY` to choose the local address
 - Also, we need to use the `htons()` function to convert the PORT NUMBER from host byte order to network byte order (recall the concept of `LITTLE-ENDIAN` and `BIG_ENDIAN` taught in CSO)



SITUATION:

There is a
problem.

Let's use
multithreading.



SOON:

SITUATION:

the
are
97
prms.oble